




Deep Sequence Modeling


Ava Amini

MIT Introduction to Deep Learning

January 6, 2025




MIT Introduction to Deep Learning
🌐 introtodeeplearning.com 🐦 [@MITDeepLearning](https://twitter.com/MITDeepLearning)



Given an image of a ball,
can you predict where it will go next?



Given an image of a ball,
can you predict where it will go next?



Given an image of a ball,
can you predict where it will go next?



Given an image of a ball,
can you predict where it will go next?



Sequences in the Wild




Audio


Sequences in the Wild



Aug 08, 2019




Sequence Modeling Applications




One to One
Binary Classification




“Will I pass this class?”
Student → Pass?




Many to One
Sentiment Classification



One to Many
Image Captioning



“A baseball player throws a ball.”




Many to Many
Machine Translation




Neurons with Recurrence


The Perceptron Revisited



Feed-Forward Networks Revisited




Feed-Forward Networks Revisited



$$x_t \in \mathbb{R}^m$$


$$\hat{y}_t \in \mathbb{R}^n$$

Handling Individual Time Steps



Neurons with Recurrence

h for
hidden
state




$$\hat{y}_t = f(\underbrace{x_t}_{\text{input}}, \underbrace{h_{t-1}}_{\text{past memory}})$$

output

Neurons with Recurrence

each stage e ekta input
& previous stage er hidden state
ta jabe as inputs




$$\hat{y}_t = f(\underbrace{x_t}_{\text{input}}, \underbrace{h_{t-1}}_{\text{past memory}})$$

output

Recurrent Neural Networks (RNNs)

Recurrent Neural Networks (RNNs)



Apply a **recurrence relation** at every time step to process a sequence:

$$\boxed{h_t} = \boxed{f_W}(\boxed{x_t}, \boxed{h_{t-1}})$$

cell state function with weights W input old state


Note: the same function and set of parameters are used at every time step

RNNs have a **state**, h_t , that is updated **at each time step** as a sequence is processed

RNN Intuition

```
my_rnn = RNN()  
hidden_state = [0, 0, 0, 0]  
  
sentence = ["I", "love", "recurrent", "neural"]
```

```
for word in sentence:  
    prediction, hidden_state = my_rnn(word, hidden_state)  
  
next_word_prediction = prediction  
# >>> "networks!"
```




RNN Intuition

```
my_rnn = RNN()
hidden_state = [0, 0, 0, 0]

sentence = ["I", "love", "recurrent", "neural"]

for word in sentence:
    prediction, hidden_state = my_rnn(word, hidden_state)

next_word_prediction = prediction
# >>> "networks!"
```




RNN Intuition

```
my_rnn = RNN()
hidden_state = [0, 0, 0, 0]


sentence = ["I", "love", "recurrent", "neural"]

for word in sentence:
    prediction, hidden_state = my_rnn(word, hidden_state)


    next_word_prediction = prediction
    # >>> "networks!"
```



RNN State Update and Output




RNN State Update and Output



Input Vector
 x_t

RNN State Update and Output



act_fun = activation
function

NN e $y = \text{act_fun}(wx+b)$ chilo

same as here,

ht o same rules folow korei update hbe

just previous hidden state add hbe &

ekthane $b = 0$

$h_t = \text{act_fun}(wx+b + w_h(t-1))$

each stage e ekta hidden state ashbe

& also ei state er ekta weight(W_{hh}) thakbe

just 2 ta multiply hoiye wx er add hoise


Update Hidden State

$$h_t = \tanh(W_{hh}^T h_{t-1} + W_{xh}^T x_t)$$

Input Vector

x_t

RNN State Update and Output



Output Vector

$$\hat{y}_t = \mathbf{W}_{hy}^T h_t$$


Update Hidden State

$$h_t = \tanh(\mathbf{W}_{hh}^T h_{t-1} + \mathbf{W}_{xh}^T x_t)$$


Input Vector

$$x_t$$

RNNs: Computational Graph Across Time



RNNs: Computational Graph Across Time



RNNs from Scratch in TensorFlow



```
class MyRNNCell(tf.keras.layers.Layer):
    def __init__(self, rnn_units, input_dim, output_dim):
        super(MyRNNCell, self).__init__()


        # Initialize weight matrices
        self.W_xh = self.add_weight([rnn_units, input_dim])
        self.W_hh = self.add_weight([rnn_units, rnn_units])
        self.W_hy = self.add_weight([output_dim, rnn_units])

        # Initialize hidden state to zeros
        self.h = tf.zeros([rnn_units, 1])

    def call(self, x):
        # Update the hidden state
        self.h = tf.math.tanh( self.W_hh * self.h + self.W_xh * x )

        # Compute the output
        output = self.W_hy * self.h

        # Return the current output and hidden state
        return output, self.h
```




RNN Implementation: TensorFlow & PyTorch


```
from tf.keras.layers import SimpleRNN  
model = SimpleRNN(rnn_units)
```




```
from torch.nn import RNN  
model = RNN(input_size, rnn_units)
```




RNNs for Sequence Modeling




One to One
"Vanilla" NN
Binary classification



Many to One
Sentiment Classification



One to Many
Text Generation
Image Captioning



Many to Many
Translation & Forecasting
Music Generation


... and many other architectures and applications



Sequence Modeling: Design Criteria

To model sequences, we need to:

1. Handle **variable-length** sequences
2. Track **long-term** dependencies
3. Maintain information about **order**
4. **Share parameters** across the sequence



Recurrent Neural Networks (RNNs) meet these sequence modeling design criteria

A Sequence Modeling Problem: Predict the Next Word

A Sequence Modeling Problem: Predict the Next Word

“This morning I took my cat for a walk.”

MIT 6.S191

A Sequence Modeling Problem: Predict the Next Word

"This morning I took my cat for a walk."

given these words

A Sequence Modeling Problem: Predict the Next Word

“This morning I took my cat for a walk.”

given these words

predict the
next word


A Sequence Modeling Problem: Predict the Next Word

“This morning I took my cat for a walk.”


given these words

predict the
next word

Representing Language to a Neural Network




Neural networks cannot interpret words




Neural networks require numerical inputs

Encoding Language for a Neural Network

word/text ke binary formate e convert korar way e hoitese one hot encoding.



Neural networks cannot interpret words



Neural networks require numerical inputs

Embedding: transform indexes into a vector of fixed size.

this cat for
my took
a | walk
morning

1. Vocabulary:
Corpus of words

a → 1
cat → 2
...
walk → N

2. Indexing:
Word to index

One-hot embedding
"cat" = $[0, 1, 0, 0, 0, 0]$
↑
i-th index

Learned embedding

A 2D plot illustrating learned word embeddings. The horizontal axis represents a semantic dimension, with 'day' and 'sun' on the left and 'happy' and 'sad' on the right. The vertical axis represents another semantic dimension, with 'run' and 'walk' on the left and 'dog' and 'cat' on the right. The words are positioned at the intersections of these axes.

3. Embedding:
Index to fixed-sized vector

Handle Variable Sequence Lengths

The food was great

vs.

We visited a restaurant for lunch

vs.

We were hungry but cleaned the house before eating

Model Long-Term Dependencies

“**France** is where I grew up, but I now live in Boston. I speak fluent ____.”



We need information from **the distant past** to accurately predict the correct word.

Capture Differences in Sequence Order



The food was good, not bad at all.

vs.


The food was bad, not good at all.



Sequence Modeling: Design Criteria

To model sequences, we need to:


1. Handle **variable-length** sequences
2. Track **long-term** dependencies
3. Maintain information about **order**
4. **Share parameters** across the sequence



Recurrent Neural Networks (RNNs) meet these sequence modeling design criteria

Backpropagation Through Time (BPTT)


Recall: Backpropagation in Feed Forward Models




Backpropagation algorithm:

1. Take the derivative (gradient) of the loss with respect to each parameter
2. Shift parameters in order to minimize loss


RNNs: Backpropagation Through Time




RNNs: Backpropagation Through Time



Standard RNN Gradient Flow




Standard RNN Gradient Flow



Computing the gradient wrt h_0 involves many factors of W_{hh} + repeated gradient computation!

Standard RNN Gradient Flow: Exploding Gradients




Computing the gradient wrt h_0 involves many factors of W_{hh} + repeated gradient computation!

Many values > 1 :
exploding gradients

Gradient clipping to
scale big gradients

Standard RNN Gradient Flow: Vanishing Gradients



Computing the gradient wrt h_0 involves many factors of W_{hh} + repeated gradient computation!

Many values > 1 :
exploding gradients

Gradient clipping to
scale big gradients


Many values < 1 :
vanishing gradients

1. Activation function
2. Weight initialization
3. Network architecture

The Problem of Long-Term Dependencies

Why are vanishing gradients a problem?

Multiply many **small numbers** together



Errors due to further back time steps
have smaller and smaller gradients




Bias parameters to capture short-term
dependencies

The Problem of Long-Term Dependencies

"The clouds are in the ____"

Why are vanishing gradients a problem?

Multiply many **small numbers** together



Errors due to further back time steps
have smaller and smaller gradients




Bias parameters to capture short-term
dependencies

The Problem of Long-Term Dependencies

Why are vanishing gradients a problem?


Multiply many **small numbers** together



Errors due to further back time steps
have smaller and smaller gradients




Bias parameters to capture short-term
dependencies



The Problem of Long-Term Dependencies

Why are vanishing gradients a problem?


Multiply many **small numbers** together



Errors due to further back time steps
have smaller and smaller gradients




Bias parameters to capture short-term
dependencies



The Problem of Long-Term Dependencies

Why are vanishing gradients a problem?


Multiply many **small numbers** together



Errors due to further back time steps
have smaller and smaller gradients




Bias parameters to capture short-term
dependencies



Gating Mechanisms in Neurons


Idea: use **gates** to selectively **add** or **remove** information within **each recurrent unit** with



Long Short Term Memory (LSTMs) networks rely on a gated cell to track information throughout many time steps.

RNN Applications & Limitations

Example Task: Music Generation



Input: sheet music


Output: next character in sheet music

Listening to
3rd movement



6.S191 Lab!

Example Task: Sentiment Classification




Input: sequence of words

Output: probability of having positive sentiment

 `loss = tf.nn.softmax_cross_entropy_with_logits(y, predicted)`

Example Task: Sentiment Classification



Tweet sentiment classification




Ivar Hagendoorn
@IvarHagendoorn

Follow



The @MIT Introduction to #DeepLearning is definitely one of the best courses of its kind currently available online introtodeeplearning.com

12:45 PM - 12 Feb 2018



Angels-Cave
@AngelsCave

Follow




Replying to @Kazuki2048

I wouldn't mind a bit of snow right now. We haven't had any in my bit of the Midlands this winter! :(

2:19 AM - 25 Jan 2019

Limitations of Recurrent Models




Limitations of RNNs



Encoding bottleneck




Slow, no parallelization



Not long memory

Goal of Sequence Modeling


RNNs: recurrence to model sequence dependencies



Goal of Sequence Modeling

RNNs: recurrence to model sequence dependencies


Limitations of RNNs




Encoding bottleneck



Slow, no parallelization



Not long memory





Goal of Sequence Modeling


Can we eliminate the need for recurrence entirely?

Desired Capabilities

 Continuous stream

 Parallelization


 Long memory




Goal of Sequence Modeling


Can we eliminate the need for recurrence entirely?

Desired Capabilities

 Continuous stream

 Parallelization

 Long memory



Goal of Sequence Modeling


Idea 1: Feed everything
into dense network

✓ No recurrence

✗ Not scalable


✗ No order

✗ No long memory



Idea: Identify and attend
to what's important

Can we eliminate the need for
recurrence entirely?



Attention Is All You Need

Intuition Behind Self-Attention


Attending to the most important parts of an input.




1. Identify which parts to attend to
2. Extract the features with high attention

Similar to a search problem!

A Simple Example: Search



Understanding Attention with Search



Query (Q)

Key (K_1)


Key (K_2)

Key (K_3)

How similar is the
key to the query?

1. **Compute attention mask:** how
similar is each key to the desired query?

Understanding Attention with Search



Query (Q)

Key (K_1)



Key (K_2)

Value (V)

Key (K_3)

2. Extract values based on attention:
Return the values highest attention

Learning Self-Attention with Neural Networks

Goal: identify and attend to most important features in input.

1. Encode **position** information
2. Extract query, key, value for search
3. Compute attention weighting
4. Extract features with high attention




Data is fed in all at once! Need to encode position information to understand order.

Learning Self-Attention with Neural Networks

Goal: identify and attend to most important features in input.

1. Encode **position** information
2. Extract query, key, value for search
3. Compute attention weighting
4. Extract features with high attention




Data is fed in all at once! Need to encode position information to understand order.

Learning Self-Attention with Neural Networks

Goal: identify and attend to most important features in input.

1. Encode **position** information
2. Extract **query, key, value** for search
3. Compute attention weighting
4. Extract features with high attention




Learning Self-Attention with Neural Networks

Goal: identify and attend to most important features in input.

1. Encode **position** information
2. Extract **query, key, value** for search
3. Compute **attention weighting**
4. Extract features with high attention

Attention score: compute pairwise similarity between each **query** and **key**

How to compute similarity between two sets of features?



Also known as the "cosine similarity"


Learning Self-Attention with Neural Networks

Goal: identify and attend to most important features in input.

1. Encode **position** information
2. Extract **query, key, value** for search
3. Compute **attention weighting**
4. Extract features with high attention

Attention score: compute pairwise similarity between each **query** and **key**

How to compute similarity between two sets of features?



Also known as the "cosine similarity"

Learning Self-Attention with Neural Networks

Goal: identify and attend to most important features in input.

1. Encode **position** information
2. Extract **query, key, value** for search
3. Compute **attention weighting**
4. Extract features with high attention

Attention weighting: where to attend to!
How similar is the key to the query?

	He	tossed	the	tennis	ball	to	serve
He	1	0	0	0	0	0	0
tossed	0	1	0	0	1	0	0
the	0	0	1	0	0	0	0
tennis	0	0	0	1	1	0	1
ball	0	1	0	1	1	0	0
to	0	0	0	0	0	1	0
serve	0	0	0	1	0	0	1

$$\text{softmax} \left(\frac{Q \cdot K^T}{\text{scaling}} \right)$$


Attention weighting

Learning Self-Attention with Neural Networks

Goal: identify and attend to most important features in input.

1. Encode **position** information
2. Extract **query, key, value** for search
3. Compute **attention weighting**
4. Extract **features with high attention**

Last step: self-attend to extract features




$$\text{softmax} \left(\frac{Q \cdot K^T}{\text{scaling}} \right) \cdot V = A(Q, K, V)$$

Learning Self-Attention with Neural Networks

Goal: identify and attend to most important features in input.

1. Encode **position** information
2. Extract **query, key, value** for search
3. Compute **attention weighting**
4. Extract **features with high attention**

These operations form a self-attention head that can plug into a larger network.
Each head attends to a different part of input.




$$\text{softmax} \left(\frac{Q \cdot K^T}{\text{scaling}} \right) \cdot V$$

Learning Self-Attention with Neural Networks

Goal: identify and attend to most important features in input.

1. Encode **position** information
2. Extract **query, key, value** for search
3. Compute **attention weighting**
4. Extract **features with high attention**




Attention is the foundational building block of the Transformer architecture.

Applying Multiple Self-Attention Heads



Attention weighting

×



Value


=




Output



Output of attention head 1



Output of attention head 2



Output of attention head 3

Self-Attention Applied

Language Processing




An armchair in the shape
of an avocado

Transformers: BERT, GPT

Devlin et al., *NAACL* 2019
Brown et al., *NeurIPS* 2020


Biological Sequences



Protein Structure Models

Jumper et al., *Nature* 2021
Lin et al., *Science* 2023

Computer Vision



Vision Transformers


Dosovitskiy et al., *ICLR* 2020



6.S191 Lab and Lectures!

Deep Learning for Sequence Modeling: Summary

1. RNNs are well suited for **sequence modeling** tasks
2. Model sequences via a **recurrence relation**
3. Training RNNs with **backpropagation through time**
4. Models for **music generation**, classification, machine translation, and more
5. Self-attention to model **sequences without recurrence**
6. Self-attention is the basis for many **large language models** – stay tuned!



6.S191: Introduction to Deep Learning

Lab 1: Deep Learning in Python and Music Generation with RNNs

Link to download labs:

<http://introtodeeplearning.com#schedule>

1. Open the lab in Google Colab
2. Start executing code blocks and filling in the #TODOs
3. Need help? Find a TA/instructor!

